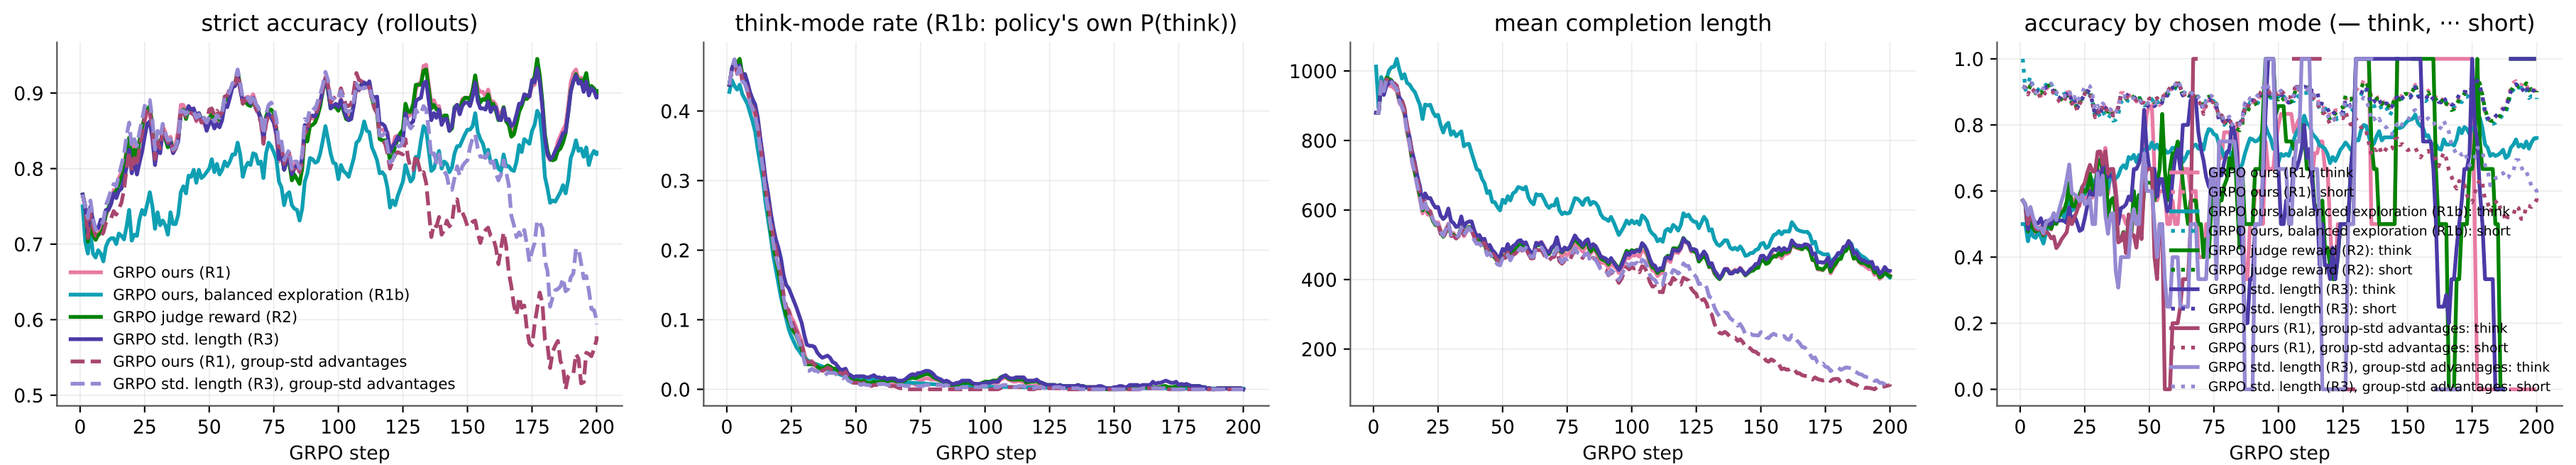
strict accuracy (rollouts)
think-mode rate (R1b: policy's own P(think))
mean completion length
accuracy by chosen mode (— think, ··· short)
GRPO step
GRPO ours (R1)
GRPO ours, balanced exploration (R1b)
GRPO judge reward (R2)
GRPO std. length (R3)
GRPO ours (R1), group-std advantages
GRPO std. length (R3), group-std advantages
GRPO ours (R1): think
GRPO ours (R1): short
GRPO ours, balanced exploration (R1b): think
GRPO ours, balanced exploration (R1b): short
GRPO judge reward (R2): think
GRPO judge reward (R2): short
GRPO std. length (R3): think
GRPO std. length (R3): short
GRPO ours (R1), group-std advantages: think
GRPO ours (R1), group-std advantages: short
GRPO std. length (R3), group-std advantages: think
GRPO std. length (R3), group-std advantages: short