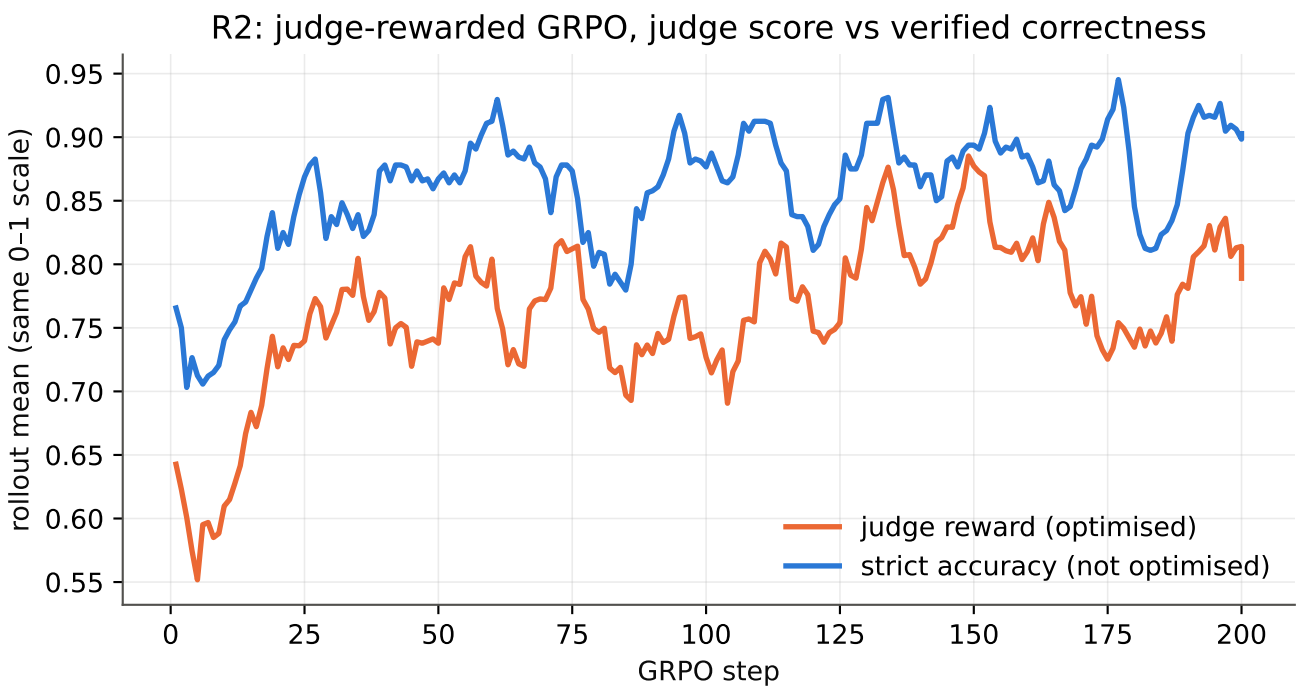
R2: judge-rewarded GRPO, judge score vs verified correctness
rollout mean (same 0–1 scale)
0.95
0.90
0.85
0.80
0.75
0.70
0.65
0.60
0.55
0
25
50
75
100
125
150
175
200
GRPO step
judge reward (optimised)
strict accuracy (not optimised)