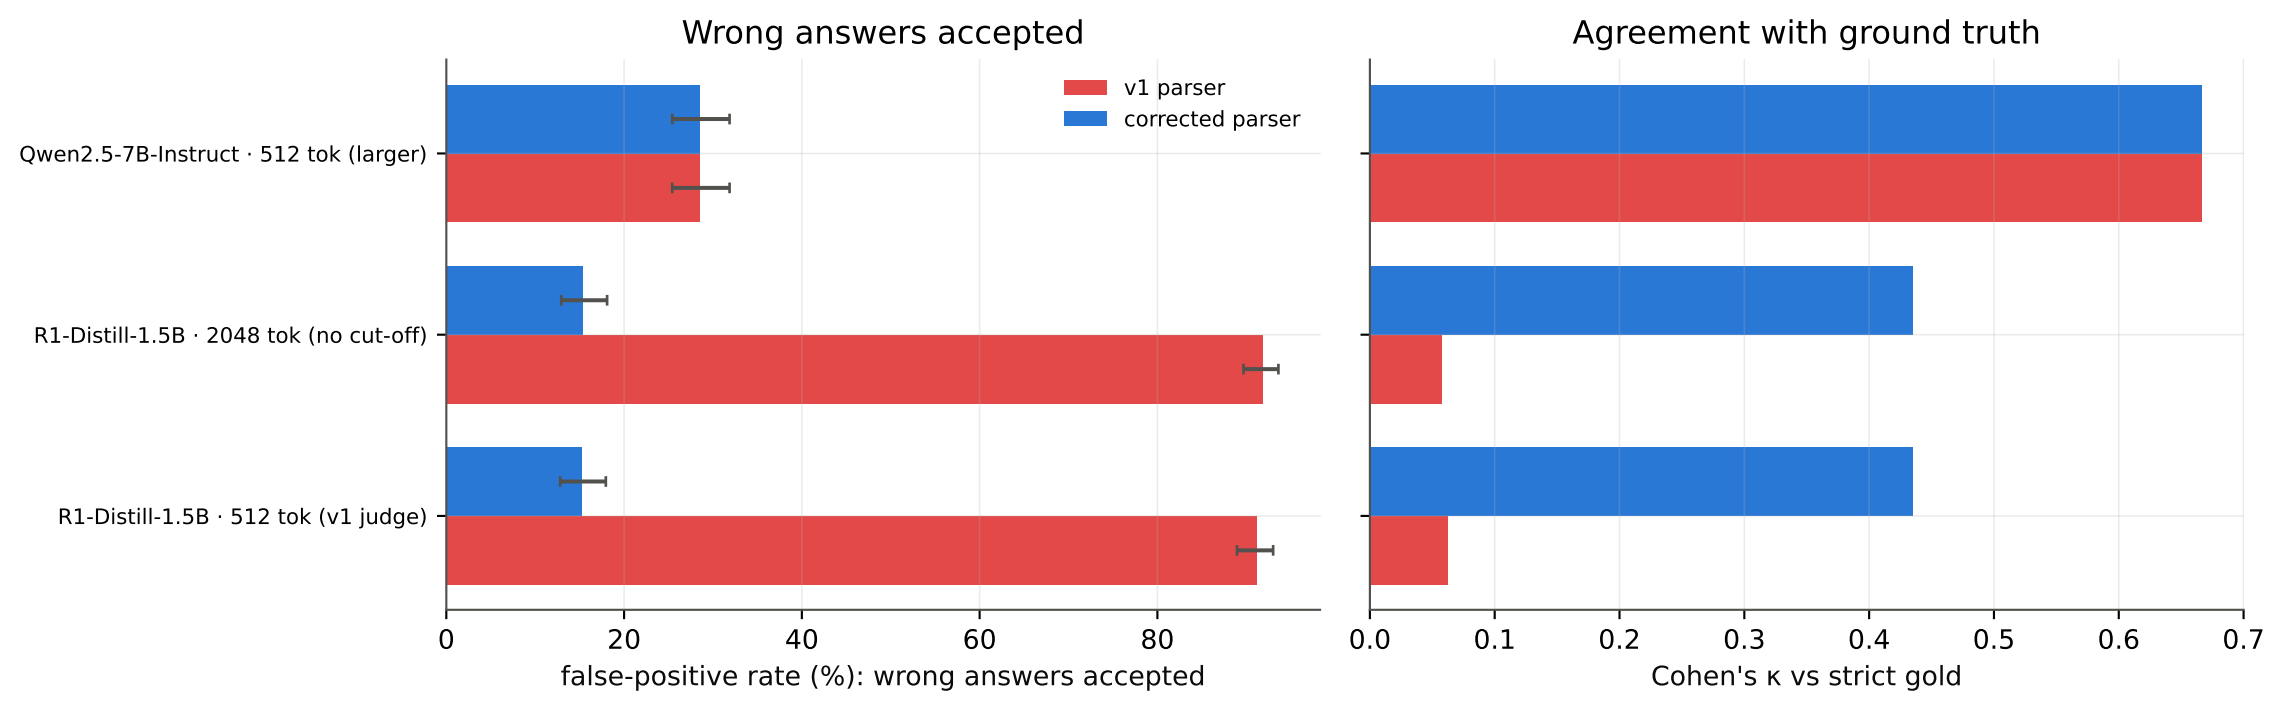
Wrong answers accepted
v1 parser
corrected parser
Qwen2.5-7B-Instruct · 512 tok (larger)
R1-Distill-1.5B · 2048 tok (no cut-off)
R1-Distill-1.5B · 512 tok (v1 judge)
0
20
40
60
80
false-positive rate (%): wrong answers accepted
Agreement with ground truth
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
Cohen's κ vs strict gold