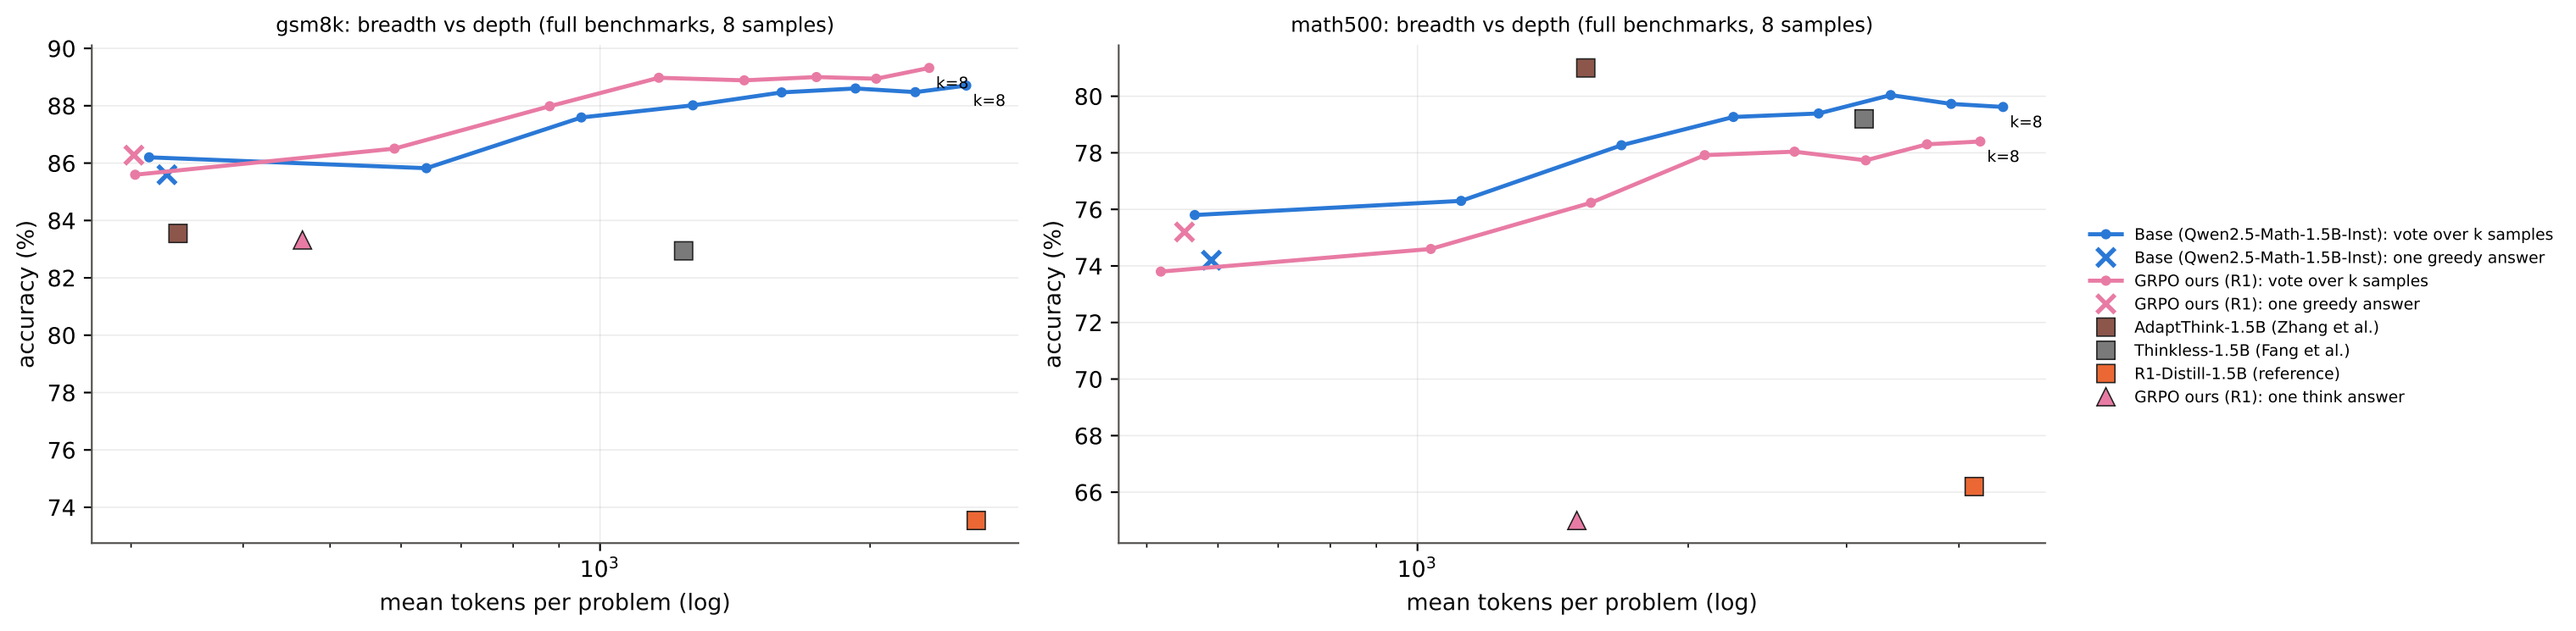
gsm8k: breadth vs depth (full benchmarks, 8 samples)
accuracy (%)
mean tokens per problem (log)
k=8
k=8
math500: breadth vs depth (full benchmarks, 8 samples)
accuracy (%)
mean tokens per problem (log)
k=8
k=8
Base (Qwen2.5-Math-1.5B-Inst): vote over k samples
Base (Qwen2.5-Math-1.5B-Inst): one greedy answer
GRPO ours (R1): vote over k samples
GRPO ours (R1): one greedy answer
AdaptThink-1.5B (Zhang et al.)
Thinkless-1.5B (Fang et al.)
R1-Distill-1.5B (reference)
GRPO ours (R1): one think answer